
GORGEO: Online Tuning for Cross-Region Network-Aware LLM Serving

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Increasingly, LLM inference services proxy client requests to engine replicas dis-
2 tributed globally. Load-balancing policies must jointly account for factors includ-
3 ing KV-cache locality, replica load, and variable network latency when optimiz-
4 ing for metrics like latency and TTFT. However, existing systems only evaluate
5 a subset of these factors in their cost model, leading to uneven concentrations of
6 load and KV-cache across replicas. We present GORGEO, a proxy architecture that
7 holistically factors network latency, prefill cost, and queueing delay using tun-
8 able parameters. Since open-source chat datasets such as LMSYS-Chat-1M and
9 WildChat-4.8M lack long-context, high prefix-reuse data, we publicly release a
10 synthetic dataset, ART-Chat-2.5M, from long-context production metadata. Intra-
11 user prefix reuse and average token length in ART-Chat-2.5M are $19\times$ and $6\times$
12 higher than in WildChat-4.8M. On a tuning window from ART-Chat-2.5M, evolu-
13 tionary strategies guide the GORGEO policy’s parameters to directly optimize p95
14 TTFT. During a held-out evaluation window, we fix the parameter values learned
15 from tuning and improve p95 TTFT by 15.5% and p95 end-to-end (E2E) latency
16 by 30.9% over baseline load-balancing policies such as simple session affinity and
17 prefix-cache.

18 1 Introduction

19 In LLM serving systems, perceived latency to the user is dominated by the time-to-first-token
20 (TTFT). On a single replica, TTFT is dominated by three costs: (i) prefill time, (ii) round trip
21 time (RTT) from client (proxy) to replica, and (iii) queueing delay behind in-flight requests. Prefix-
22 caching, which is enabled in inference engines SGLang [Zheng et al., 2024b] and vLLM [Kwon
23 et al., 2023], eliminates the prefill cost of previous turns in a multi-turn conversation. As LLM con-
24 text windows increase in length, the time saved by prefix caching 90% of a prompt with 100,000
25 tokens reduces the prefill cost to 10,000 tokens and decreases TTFT substantially.

26 Since LLM deployments proxy requests to inference engines across regions, the cost savings of
27 prefix-caching depend on choosing a replica with the request session’s prefix. Popular routing poli-
28 cies such as consistent hashing and prefix reuse aim to distribute load evenly while creating affinity
29 between a session’s requests and replica(s) to maximize KV-cache reuse [Karger et al., 1997, Stoica
30 et al., 2003, Zheng et al., 2024b, Xia et al., 2025]. In compute-constrained regimes, bursty work-
31 loads can saturate a high-affinity replica, causing head-of-line (HOL) queueing delays from decode
32 memory contention and negating cost savings from prefix-cache reuse. Routing policies that holis-
33 tically evaluate all costs related to TTFT can maintain high prefix-cache reuse while minimizing the
34 negative effects of load saturation and heterogeneous network latency.

35 Existing load balancing policies such as least-load, session affinity, and prefix-reuse may account for
36 replica load or KV-cache hit rate; however, no existing policies consider network latency in cross-

region scenarios, which can range on the order of 10ms to 1s (Appendix A). GORG0’s routing policy accounts for all three TTFT costs, normalizes the units of measurement via tunable parameters, and jointly optimizes parameters through online tuning on real user workloads. To tune and stress-test different routing policies, in (§3) we compile an LLM traffic trace from real production requests with high prefix-reuse and long-context prompts. The trace follows Mooncake’s FAST’25 format [Qin et al., 2025] containing per-request timestamps, which can be linearly scaled to simulate different load profiles for various hardware. On a held-out window, fixed-weight GORG0 improves p95 TTFT by 15.5% and p95 E2E latency by 30.9% over session affinity (§4).

2 GORG0

2.1 Design space

For LLM load balancing policies, the design space includes which routing signals to consider and methods for choosing environment settings.

Routing signals: Across existing load balancing policies, two distinct routing signals are considered: (i) replica load and (ii) KV-cache locality. Simple policies such as least-load and least-request aim to evenly balance replica load while prefix-cache and simple-session-affinity aim to maximize KV-cache reuse per request. NVIDIA Dynamo [NVIDIA, 2026b], SGLang Router [Zheng et al., 2024b], and Preble [Srivatsa et al., 2025] offer algorithms that jointly optimize for replica load and KV-cache locality.

GORG0 takes inspiration from early content-delivery network design, which has many similarities to LLM serving. Wang et al. [2002] used three signals—server load, network proximity, and cache locality—to reduce response time and increase capacity. Evaluating network proximity can improve responsiveness, or TTFT, in geo-distributed serving environments. GORG0 is the first LLM load balancing policy to also optimize (iii) network latency from client to server.

Weight-selection methods: Methods for tuning the weights of a request scheduler fall into (i) offline and (ii) online protocols. Existing offline methods include Dynamo’s SLA Planner, which has a sophisticated method to autoscale prefill and decode workers to meet TTFT and ITL targets. It requires profiling before deployment to determine the optimal tensor parallelism configurations and scaling parameters for the planner to use during runtime [NVIDIA, 2026a].

GORG0 uses online calibration because cache state, queue state, and batching behavior are intrinsically influenced by the routing policy’s decisions, and accurately replicating the conditions necessary to evaluate the policies requires a dedicated replay fleet, shadow execution, or a sufficiently accurate simulator of cache and continuous-batching dynamics. After all, a request log generated under one policy does not contain the counterfactual cache and queue evolution that would have arisen under another set of weights.

2.2 Cost model

TTFT is known to consist of three different costs: network latency, prefill time, and queueing delay [He et al., 2025]. The KV-Cache, which stores key-value pairs of input sequences, removes redundant prefill computation for user sequences sharing a prefix with previously computed sequences [Zheng et al., 2024b]. In a cross-region deployment setting, for an inference engine replica i holding a set of cached token prefixes c_i , the cost of TTFT can be defined as the following function, where x_r is the input sequence of tokens for a request r and the prefill time depends only on the set difference $(x_r \setminus c_i)$, the tokens in x_r not already cached on i :

$$\text{TTFT} = T_{\text{network}}(i) + T_{\text{queue}}(i) + T_{\text{prefill}}(x_r \setminus c_i) \quad (1)$$

Theoretically, T_{network} and T_{queue} correspond with round trip time from a client to server and the duration of request processing ahead of the incoming request r . However, inference engines support batching requests continuously [Yu et al., 2022], and admission is still bounded by a maximum number of concurrent requests and the available KV-cache budget [Kwon et al., 2023].

In a distributed system, the temporal delay of retrieving queueing metrics from an engine replica makes cost evaluation challenging. We represent the input to T_{queue} as the total number of tokens

Table 1: Dataset. ART-Chat-2.5M contains higher average input tokens and global prefix reuse, which is measured by adding the intra-user reuse and cross-user reuse, over other chat datasets.

Dataset	Total reqs	Users	Avg input tokens	Intra-user reuse	Global reuse
LMSYS-Chat-1M	1,000,000	—	467	—	3.4%
WildChat-4.8M	3,199,860	1,833,730	2,925	4.7%	32.5%
ART-Chat-2.5M	2,525,215	4,984	17,964	89.4%	89.7%

of requests without a completion event on the client proxy. T_{network} is measured trivially as the exponential weighted moving average of a ping’s round trip time from client proxy to server.

$$\text{TTFT} = T_{\text{network}}(\text{RTT}_i) + T_{\text{queue}} \left(i, \sum_{j=1}^{n_r} x_j \right) + T_{\text{prefill}}(x_r \setminus c_i) \quad (2)$$

2.3 GORGO proxy design

The GORGO proxy routes client requests to the engine replica with the minimum calculated cost from Equation 2, parameterized by weights W_{rtt} , W_{prefill} , and W_{queue} . These parameters weight the inputs T_{network} , T_{queue} , and T_{prefill} , which unfairly mixes the units of time and tokens, to normalize the correlated costs of latency, prefill cost, and queueing time on TTFT. The weight of T_{prefill} is fixed to 1 because GORGO proxy makes routing decisions based on relative replica cost: only the ratio of weights matters in this design.

$$\text{TTFT} = W_{\text{rtt}} * T_{\text{network}}(\text{RTT}_i) + W_{\text{queue}} * T_{\text{queue}} \left(i, \sum_{j=1}^{n_r} x_j \right) + T_{\text{prefill}}(x_r \setminus c_i) \quad (3)$$

GORGO proxy uses a simple (1+1) evolutionary strategy to tune weights W_{rtt} and W_{queue} on the objective function, p95 TTFT. Each parent weight $x_{t,k}$ is perturbed multiplicatively in log-space by a normal random variable z_k times step size σ , and the new offspring weight x'_k is clamped to values in the hyperparameter range $[lo_k, hi_k]$ where $lo_k > 0$ and $hi_k > 0$.

$$x'_k = \text{clip}(\exp(\ln(x_{t,k}) + \sigma_t * z_k), [lo_k, hi_k]) \quad (4)$$

When x' beats the parent weight x_t on the objective metric, the incumbent weight x_{t+1} is updated to x' , and σ is adjusted to maintain Rechenberg’s 1/5 success rule [Rechenberg, 1973] of roughly one accepted offspring for every five proposals.

3 Dataset

Existing LLM chatbot datasets lack two critical components for benchmarking cache-aware policies: (i) prefill-bound requests with long-context prompts and (ii) multi-turn workloads with high prefix-reuse between requests. For example, we measure the average request length and global prefix reuse of **LMSYS-Chat-1M** [Zheng et al., 2024a] and **WildChat-4.8M** [Zhao et al., 2024], two popular LLM datasets derived from public chatbot demos (Table 1).

ART-Chat-2.5M is a long-context, multi-turn dataset synthetically generated from a week-long metadata trace of production inference traffic with the same prefix-reuse structure as the original workload. We release a replay-ready trace in the Mooncake FAST’25 format, which contains per-request timestamps, request metadata, and synthetically generated chat completion data [Qin et al., 2025]. The dataset is included in the anonymous supplementary material. By storing request timestamps, one can linearly scale the time between requests to simulate different load profiles for various hardware. Notably, the intra-user prefix reuse and average token length in ART-Chat-2.5M are $19\times$ and $6\times$ higher than in WildChat-4.8M. Dataset analysis is in Appendix B.

4 Results

Results across two windows. We evaluate on a three-region fleet of L40S GPUs running Qwen3.5-35B-A3B in the SGLang inference engine (baselines and windows in Appendix A). GORGO tunes w_{rtt} and w_{queue} on Apr 5 to minimize rolling p95 TTFT, then freezes the learned weights ($w_{\text{rtt}}=0.276$, $w_{\text{queue}}=0.5$) on Apr 6 (Apr 7 in Appendix E).

Table 2: Experiment results. In the tuning window, GORGO learns weights $w_{\text{rtt}}=0.276$, $w_{\text{queue}}=0.5$ after initializing at weights $w_{\text{rtt}}=0.5$, $w_{\text{queue}}=0.1$. On the held-out evaluation window, GORGO’s weights are frozen, and the policy outperforms every baseline policy on p95 TTFT.

Policy	TTFT p50 (ms)	TTFT p95	TTFT p99	E2E p50 (s)	E2E p95	E2E p99	ITL (ms)
<i>Tuning window — Apr 5 16:15–16:45 (7,195 requests)</i>							
GORGO	673	2,514	4,473	3.04	8.91	17.54	17.6
simple-session-affinity	712	2,428	5,190	3.48	18.27	22.61	20.2
least-load	763	2,447	4,349	3.63	13.69	17.57	21.7
least-request	968	3,970	7,012	4.37	16.62	20.81	25.2
prefix-cache	1,477	6,784	9,613	13.14	22.63	26.81	82.3
<i>Held-out eval — Apr 6 15:05–15:35 (7,630 requests)</i>							
GORGO	491	1,584	2,222	1.80	3.28	4.37	9.0
simple-session-affinity	627	1,875	3,056	2.01	4.75	7.34	10.3
least-load	616	1,818	2,885	2.17	4.06	5.78	10.7
least-request	634	1,852	2,484	2.08	3.99	5.23	9.9
prefix-cache	574	1,798	2,750	2.07	4.95	10.34	10.6

While the GORGO policy’s weights are updated to minimize p95 TTFT during tuning, the held-out, fixed-weight evaluation shows generalization of the learned values across days, with GORGO improving p95 TTFT by 15.5% and E2E latency by 30.9% over session-affinity. GORGO slightly underperforms baseline policies on the tuning window because the evolutionary strategy actively explores the space of parameters and tests worse weights than the learned solution, which converges after 672 samples (Appendix D).

BFCL multi-turn replay. As an additional high-prefix-reuse workload, we replay 200 teacher-forced episodes from BFCL’s `multi_turn_base` subset, where each episode contains sequential tool calls [Patil et al., 2025]. We use two Qwen3-4B H200 replicas in us-east and us-west, 128 concurrent episodes, and at most 256 output tokens per request.

Table 3: BFCL multi-turn replay. Bold marks the best value per column.

Policy	TTFT p50 (ms)	TTFT p95	TTFT p99	E2E p50 (s)	E2E p95	E2E p99	ITL (ms)
GORGO	3,328	4,700	5,720	3.66	5.17	6.15	—
simple-session-affinity	3,655	5,087	5,608	4.11	5.61	6.17	—

Under high load, GORGO reduces p50 and p95 TTFT by 8.9% and 7.6%, and p50 and p95 E2E latency by 10.8% and 7.9%, respectively.

Exploiting continuous batching in SGLang. In the continuous batching paradigm, GORGO can stumble upon abnormal weight values and achieve an impressively low p95 TTFT at the cost of high p95 E2E latency. Appendix G presents results from an experiment where the evolutionary strategy learns to set w_{queue} to 0 while keeping w_{rtt} near 0.23, which results in p95 TTFT 17% better than the next-best policy. For this window, GORGO chose to send 100% of requests to the closest replica in us-ashburn. Once the request enters the memory-bound decode phase, the replica’s memory headroom saturates, KV-cache thrashes between GPU and CPU memory, and ITL/E2E latency collapses [Yu et al., 2022, Kwon et al., 2023]. To mitigate this, we set a non-zero lower bound for the load-term, which forces the policy to balance request concentration across replicas.

References

- Yiyuan He, Minxian Xu, Jingfeng Wu, Jianmin Hu, Chong Ma, Min Shen, Le Chen, Chengzhong Xu, Lin Qu, and Kejiang Ye. BanaServe: Unified KV cache and dynamic module migration for balancing disaggregated LLM serving in AI infrastructure, 2025.
- David Karger, Eric Lehman, Tom Leighton, Rina Panigrahy, Matthew Levine, and Daniel Lewin. Consistent hashing and random trees: Distributed caching protocols for relieving hot spots on the world wide web. In *Proceedings of the Twenty-Ninth Annual ACM Symposium on Theory of Computing (STOC)*, pages 654–663, 1997.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*, 2023.
- NVIDIA. NVIDIA Dynamo: Profiler. <https://docs.nvidia.com/dynamo/components/profiler>, 2026a. Accessed July 28, 2026.
- NVIDIA. NVIDIA Dynamo: Router guide. <https://docs.nvidia.com/dynamo/latest/user-guides/kv-cache-aware-routing>, 2026b. Accessed July 28, 2026.
- Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. The berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, volume 267, pages 48371–48392. PMLR, 2025.
- Ruoyu Qin, Zheming Li, Weiran He, Mingxing Zhang, Yongwei Wu, Weimin Zheng, and Xinran Xu. Mooncake: A KVCache-centric disaggregated architecture for LLM serving. In *23rd USENIX Conference on File and Storage Technologies (FAST)*, 2025.
- Qwen Team. Qwen3 technical report, 2025.
- Ingo Rechenberg. *Evolutionsstrategie: Optimierung technischer Systeme nach Prinzipien der biologischen Evolution*. Frommann-Holzboog, Stuttgart, 1973.
- Vikranth Srivatsa, Zijian He, Reyna Abhyankar, Dongming Li, and Yiyang Zhang. Preble: Efficient distributed prompt scheduling for LLM serving. In *International Conference on Learning Representations (ICLR)*, 2025.
- Ion Stoica, Robert Morris, David Liben-Nowell, David R. Karger, M. Frans Kaashoek, Frank Dabek, and Hari Balakrishnan. Chord: A scalable peer-to-peer lookup protocol for internet applications. *IEEE/ACM Transactions on Networking*, 11(1):17–32, 2003.
- The AIBrix Team. AIBrix: Towards scalable, cost-effective large language model inference infrastructure, 2025.
- Limin Wang, Vivek S. Pai, and Larry L. Peterson. The effectiveness of request redirection on CDN robustness. In *5th Symposium on Operating Systems Design and Implementation (OSDI)*. USENIX Association, 2002. doi: 10.1145/844128.844160.
- Tian Xia, Ziming Mao, Jamison Kerney, Ethan J. Jackson, Zhifei Li, Jiarong Xing, Scott Shenker, and Ion Stoica. SkyWalker: A locality-aware cross-region load balancer for LLM inference, 2025.
- Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. Orca: A distributed serving system for Transformer-based generative models. In *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 2022.
- Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, and Yuntian Deng. WildChat: 1M ChatGPT interaction logs in the wild. In *International Conference on Learning Representations (ICLR)*, 2024.

- 186 Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Tianle Li, Siyuan Zhuang, Zhanghao Wu, Yong-
187 hao Zhuang, Zhuohan Li, Zi Lin, Eric P. Xing, Joseph E. Gonzalez, Ion Stoica, and Hao Zhang.
188 LMSYS-Chat-1M: A large-scale real-world LLM conversation dataset. In *International Confer-*
189 *ence on Learning Representations (ICLR)*, 2024a.
- 190 Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao,
191 Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. SGLang:
192 Efficient execution of structured language model programs. In *Advances in Neural Information*
193 *Processing Systems (NeurIPS)*, 2024b.

194 A Experimental setup

195 **Proxy and engine configuration.** Each policy runs the GORGO proxy on a small CPU worker
 196 in us-ashburn and controls a dedicated SGLang inference engine in each of the following regions:
 197 us-ashburn, eu-frankfurt, and ap-seoul. The engines contain two L40S GPUs each and serve the
 198 Qwen3.5-35B-A3B model in FP8 format [Qwen Team, 2025]. All policies are benchmarked on
 199 the same workload in parallel to rule out any variance in network conditions. The round-trip time
 200 between the us-ashburn proxy and engines during the tuning window is plotted in Figure 1. Due to
 201 the Qwen model’s limited context length of 32,768 tokens, we filter out any requests that contain
 202 $>24,000$ tokens to leave adequate KV headroom. In SGLang, we set `max_concurrent_requests`
 203 to 64 and `max_output_tokens` to 128 to limit unnecessary decode while simulating adequate load
 204 on the replica. All workloads run alongside the proxy, dispatching requests to a local chat completion
 205 endpoint.




Figure 1: Round trip time (RTT) between the GORGO policy’s proxy in us-ashburn and the SGLang engines in us-ashburn, eu-frankfurt, and ap-seoul over the Apr 5 16:15–16:45 tuning window. The proxy measures RTT every 30 seconds on a separate probe from the `/metrics` request and smoothes the RTT with an exponential moving average (EWMA). Each region’s mean latency demonstrates that network latency either adds a consideration when choosing between replicas with disparate load and prefill cost or simplifies the choice when choosing between replicas with similar costs.

206 **Baseline policies.** We benchmark the GORGO policy and compare performance of both online
 207 and static modes to the below baselines. All SGLang metrics are scraped every 30 seconds from the
 208 engine’s Prometheus `/metrics` endpoint.

- 209 1. `least-load` minimizes the sum of proxy-tracked queued requests with SGLang met-
 210 rics `num_running_reqs`, `num_queue_reqs`, and `num_used_tokens`. Queued requests
 211 to the proxy are defined as recently dispatched requests without a token response, and
 212 `num_used_tokens` are the currently occupied per-token KV slots.
- 213 2. `least-request` chooses the replica with the fewest in-flight requests from the proxy.
- 214 3. `prefix-cache` matches the request’s prefix to the replica with the highest prefix-cache
 215 overlap, tracked on the proxy-side by a prefix trie of dispatched requests [The AIBrix Team,
 216 2025].
- 217 4. `simple-session-affinity` hashes the first 256 tokens from a request and routes to the
 218 replica with that prefix hash.

219 **Tuning and evaluation windows.** We pick three 30-minute windows with high user diversity
 220 from the ART-Chat-2.5M trace: Apr 5th 16:15–16:45, Apr 6 15:05–15:35, and Apr 7 19:45–20:15.
 221 Statistics on each of these windows are found in Table 4. We assign Apr 5th as the window where
 222 we tune GORGO’s weights online to minimize the p95 TTFT of a rolling 128-request window with
 223 hop size 32. w_{rtt} and w_{queue} are each initialized to 0.5 and 0.1 and restricted to ranges $[0.05, 2.0]$
 224 and $[0.05, 0.5]$. These values are hand-picked from the paradigm of continuous batching in SGLang,

where incoming requests can be scheduled into the current batch, affecting TTFT less significantly than a fixed floor of network latency between regions. The Apr 6–7 windows fix weight values GORGO learned on the Apr 5 tuning window.

B Dataset analysis

WildChat-4.8M contains hashed IPs per request, allowing categorization of cross-user and intra-user reuse while LMSYS-Chat-1M lacks user identification. Additional results from benchmarking GORGO on WildChat-4.8M can be found in Appendix F. Figure 2 compares the three traces. ART-Chat-2.5M’s reuse is almost entirely intra-user (89.4%) with 0.3% cross-user reuse, matching long-context multi-turn sessions. WildChat’s reuse is predominantly cross-user (28%), from shared templates and a common system prompt at a much shorter mean length (2,925 tokens). LMSYS has 3.4% global reuse and no measurable intra-user reuse.




Figure 2: *Left*: Dataset characterization. ART-Chat-2.5M’s prefix reuse is overwhelmingly intra-user (89%) with minimal cross-user reuse (0.3%), which reflects the long-context, multi-turn data; WildChat’s prefix reuse is predominantly cross-user (28%, shared templates); LMSYS has minimal cross-user reuse (3%) and no intra-user reuse due to the lack of a field for client origin. *Right*: ART-Chat-2.5M contains the maximum for each attribute except for cross-user reuse, which in WildChat is a function of the shorter average token length (2,925 tokens) and a common system prompt.

C Evaluation window characteristics

Table 4: Workload characteristics of the tuning and held-out windows. Apr 7 is reported in Appendix E.

	Apr5 tuning 16:15–16:45	Apr6 eval 15:05–15:35	Apr7 eval 19:45–20:15
Distinct users	579	606	540
Requests / user	12.8	12.9	16.4
Avg turns / conv.	30.3	27.7	21.5
Multi-turn requests	73.9%	74.1%	73.7%
Median input tokens	5,638	5,787	5,745
Avg input tokens	6,877	7,008	7,148
p95 input tokens	19,886	20,974	20,435
Block global reuse	78.8%	77.6%	87.9%
Block intra-user reuse	78.5%	77.1%	87.4%

D Evolutionary strategy convergence




Figure 3: Convergence of the evolutionary strategy on GORGO policy weights w_{queue} and w_{rtt} . Both weights reach their local optima after 672 samples, which occurs on the 18th evolutionary step. The fitness function is the negative p95 TTFT (in seconds) of the rolling 128-request tuning window, so a smaller magnitude is better; the best score reached is -1.276 , i.e. a 1.276 s best-window p95. Because it is computed over the rolling tuning window, this is lower than the 2,514 ms overall tuning-window p95 in Table 2.

238 E Apr 7 held-out window

Table 5: Apr 7 19:45–20:15 held-out window.

Policy	TTFT p50 (ms)	TTFT p95	TTFT p99	E2E p50 (s)	E2E p95	E2E p99	ITL (ms)
GORGO	398	1,417	2,061	1.74	3.36	6.81	9.3
simple-session-affinity	457	1,522	2,402	1.74	3.92	7.12	9.3
least-load	495	1,669	2,368	1.95	4.07	6.22	10.0
least-request	527	1,759	2,578	1.98	4.44	7.85	10.0
prefix-cache	533	1,861	2,872	2.20	12.00	20.04	11.8

239 F WildChat replay

Table 6: Held-out WildChat-4.8M replay (5.7 requests/s open-loop), $c=64$, using one H100 replica in each of us-west4, CANADA-2, and sines-2. Fixed-weight GORGO uses weights learned on the first window ($w_{\text{rtt}}=2.0$, $w_{\text{queue}}=0.104$). Each arm contains 20,422 completed requests with zero failures. Bold marks the best value per column; ITL is the median.

Policy	TTFT p50 (ms)	TTFT p95	TTFT p99	E2E p50 (s)	E2E p95	E2E p99	ITL (ms)
GORGO	215	506	728	1.55	2.42	2.98	10.1
random	244	611	1,166	1.38	2.51	4.04	8.4

240 G Load weight adaptation

241 Table 7 and Figure 4 give the Apr 2 window where $w_{\text{queue}}=0$ routes $\sim 100\%$ of traffic to one replica.
 242 Tuned gorgo-static wins every TTFT percentile but is worst on E2E and ITL tails.

Table 7: Reward hacking on the Apr 2 12:30–13:00 window. With $w_{\text{queue}}=0$, tuned `gorgo-static` wins every TTFT percentile but is worst on E2E and ITL tails, because it routes $\sim 100\%$ of traffic to one replica. Bold marks the best value per column; ITL is the median.

Policy	TTFT p50 (ms)	TTFT p95	TTFT p99	E2E p50 (s)	E2E p95	E2E p99	ITL (ms)
<i>Held-out eval — Apr 2 12:30–13:00 (3,071 requests)</i>							
<code>gorgo-static</code> (hacked)	180	1,072	1,810	2.05	12.49	16.63	14.1
<code>simple-session-affinity</code>	268	1,715	2,947	1.74	3.88	5.90	10.8
<code>least-load</code>	278	1,669	2,665	1.69	3.89	8.94	10.1
<code>least-request</code>	274	1,285	1,902	1.37	2.61	3.34	8.6
<code>random</code>	283	1,456	2,124	1.40	2.92	4.24	8.5
<code>prefix-cache</code>	288	1,468	2,374	1.56	3.59	7.12	9.4




Figure 4: *Left*: TTFT p95 (dark) vs. E2E p95 (light) on the midday diurnal trace with $w_{\text{queue}}=0$. `gorgo-static` wins TTFT but its E2E inflates to 12.6 s. *Right*: routing concentration. `gorgo-static` sends 100% of requests to one replica.